

Build for Billions

Initial Screening Proposal

CasePilot AI — A source-linked legal case preparation assistant

Team Name: MIB

Team Leader: Sharad Siddanagoudar

Team Members: Akshata Patil, Aayush Kumar Rai

College/University: KLE Technological University

Selected Track: Agentic AI

1. Problem Statement

Small and mid-sized law firms spend significant time manually reviewing documents, reconstructing case histories, finding missing information, and organizing follow-up work. Facts are scattered across agreements, payment records, notices, emails, and other evidence. This makes case preparation slow and makes it difficult to verify where a claim came from.

Independent lawyers and small legal teams may have limited staff to perform this work. In a property dispute, inconsistent possession dates or payment claims can be overlooked when documents are reviewed separately. The challenge is to turn a case file into a structured, verifiable brief while keeping legal judgment with the lawyer.

2. Proposed Solution

CasePilot AI is a proposed AI case-preparation assistant that converts uploaded documents into an evidence-linked case brief. A controlled agent workflow extracts information, compares accounts across documents, identifies potential gaps, and proposes follow-up tasks for lawyer review.

Users create a case and upload text-based PDF, DOCX, or TXT files. The system preserves page or passage references during extraction, identifies parties and factual claims, and generates a summary and timeline. Lawyers can inspect the evidence behind findings and edit, approve, or reject suggested tasks.

Source-linked brief. Key parties, facts, and events link to supporting passages; PDF page references are shown where available.

Conflicts and gaps. The system presents competing accounts and flags information not found in the uploaded material for human verification.

Review and history. Suggested tasks remain pending until reviewed. An activity history records AI suggestions and user decisions.

The 24-hour prototype will demonstrate a fictional property possession and payment dispute using approximately six short documents. Scanned files will be flagged as needing OCR. CasePilot supports preparation; it does not make autonomous legal decisions or execute consequential legal actions without human approval.

3. Target Users and Impact

Primary users are independent lawyers, small and mid-sized law firms, paralegals, and legal-aid teams. The intended benefit is less repetitive document organization and faster access to evidence, helping limited legal staff prepare cases more consistently. Clients in underserved communities could benefit indirectly when legal-aid teams can use their time more effectively. These benefits are proposed and will require user validation.

4. Technology Stack

Frontend: React, TypeScript, and Tailwind CSS for uploads, the case workspace, source viewing, and task review.

Backend: Node.js with TypeScript for document parsing, job processing, AI orchestration, and API endpoints.

Data and access: Supabase Postgres, private Supabase Storage, and Supabase Auth for case records, files, and user access. Case ownership checks and row-level security will restrict access.

AI integration: A hosted language model through the NVIDIA API Catalog, subject to available developer access and quota. The final model will be selected after testing extraction quality, latency, and structured output. API keys remain on the backend.

Processing: PDF text extraction and DOCX/TXT parsing preserve source identifiers. JSON schema validation and citation checks validate output structure and source references; legal interpretation remains subject to review.

Planned flow: Upload → extract text → extract evidence → analyze the case → validate references → lawyer review. Processing status is stored in Postgres and shown in the interface. The prototype will analyze a bounded set of short documents.

5. Innovation and Scalability

The proposed differentiation is an evidence-first review workflow: summaries, timeline events, conflicts, and tasks share traceable source records. Conflicting accounts remain visible rather than being silently resolved. Document statements are distinguished from AI inferences, and missing material is described as absent from the uploaded file set, not proven nonexistent.

The agentic workflow connects document understanding to proposed follow-up work, with an explicit human approval boundary. This allows lawyers to verify a finding, inspect its source, and decide what to do next in one workspace. The prototype will use fictional data and will not include external messaging or legal research.

Beyond the hackathon, processing could move to dedicated workers with queued jobs, document chunking, and retrieval for larger files. Future extensions could include OCR, regional-language support, and workflows for tenancy, consumer, and contract disputes. Broader deployment would require stronger tenant isolation, retention controls, security evaluation, and testing with legal professionals.

6. Prototype and References

Current status: proposal and concept wireframe; a working prototype is planned for the hackathon. The wireframe illustrates document intake and evidence review. All displayed case facts and excerpts are fictional.

Wireframe reference: Appendix A in this document and the accompanying CasePilot_Wireframe.png file.

Planned AI platform: <https://build.nvidia.com/>

NVIDIA API integration guide: <https://docs.api.nvidia.com/nim/docs/api-quickstart>

Supabase documentation: <https://supabase.com/docs>

7. Team Contribution

Sharad Siddanagoudar — Team Leader: System architecture, backend development, Supabase integration, document processing, and AI orchestration.

Akshata Patil: Frontend development for case creation, uploads, and source viewing; wireframe refinement and submission documentation.

Aayush Kumar Rai: Frontend development for the brief, timeline, issue review, and task decisions; documentation and demo preparation.

The team will jointly test and demonstrate the complete workflow.

Appendix A Concept Wireframe

Two proposed screens show the path from document upload to evidence inspection and lawyer review. This is a design reference, not a functioning application.

CasePilot AI

Concept wireframe | Fictional property dispute | No live processing

01 Create a case and upload evidence

CASEPILOT / New case

Case title

Rao v Mehta - Possession and payment

Case type: Property dispute

Drop documents here or browse files

Text-based PDF, DOCX and TXT

Scanned documents are flagged for OCR

Documents added

Sale agreement.pdf

Ready

Seller notice.pdf

Ready

Buyer email.docx

Ready

Analyze documents

Progress: extracting text > identifying evidence > preparing review

02 Review findings and inspect their sources

CASEPILOT / Rao v Mehta

AI draft - lawyer review

DOCUMENTS

Sale agreement

Payment receipt

Bank record

Seller notice

Buyer email

Property annexure

Brief

Timeline

Issues

Tasks

Possible conflict

Possession handover

Seller states keys were delivered on 10 June.

Seller notice - p. 2

Buyer states keys had not arrived by 15 June.

Buyer email - passage 3

Information gap

No signed handover acknowledgment was found in the uploaded material.

Activity history

AI suggested follow-up > Lawyer decision recorded

SOURCE EVIDENCE

Seller notice.pdf / Page 2

"Possession of the property and the keys were handed over to the buyer on 10 June 2026."

Proposed follow-up task

Request the signed handover acknowledgment.

Reason: conflicting possession accounts.

Approve

Edit

Reject

Status: Proposed - awaiting lawyer review

Approval records a task decision; it does not send a request externally.

Reference only. Source excerpts and case details are fictional.




Figure 1. CasePilot AI — Conceptual Dashboard and Interface Mockups. AI-generated visual created to illustrate the proposed dashboard, case workspace, document review, timeline, and task-management features. The image is for conceptual demonstration only and does not represent a functional prototype or final UI.

Source: Generated using OpenAI ChatGPT (AI image generation), September 2026.

Demonstration sequence: upload the fictional documents, inspect a possession-date conflict, open both source references, and record a decision on the proposed follow-up task.